

Requirements vs. Implementation — Progress Update

What changed since this morning's gap analysis, measured the same way, against the same briefing and the same 26-requirement scorecard — so it reads as a direct continuation, not a new report.

BASELINE DOCUMENT	"Innovation Radar V2 — Requirements vs. Implementation", sent this morning · 26 requirements assessed, 7 MET / 6 PARTIAL / 8 MISSING / 5 NOT REQUIRED
SCOPE	radar_v2/ (Reflex app) + Pipelineteamfile/ (data pipeline) — unchanged
METHOD	Direct source reading; live team database queried read-only; every claim below re-verified against the running app before writing it down
LIVE DATA	21 real opportunity spaces, 41 articles, 34 classifications — same corpus as this morning, now scored and typed differently
WORK PRODUCT	207 automated tests passing · 54 changed or new files · nothing committed yet
GENERATED	25 August 2026, end of day

Status labels — same legend as the baseline document

MET

Built and working as the briefing describes.

PARTIAL

Something exists, but not in the shape the briefing asks for.

MISSING

A stated requirement with no implementation.

NOT REQUIRED


The briefing itself marks it next-phase, illustrative or optional.

Executive summary

Where the application stands now, against the exact gaps the baseline document flagged.

The baseline report's single biggest gap — **"the scoring model does not explain itself, and the number it produces is not an attractiveness score"** — is closed. The old "Strategic fit" figure (article count ÷ largest space's article count) is gone. In its place is a five-component, weighted Attractiveness score with every component visible, and a fully independent Now/Next/Later timing rule that no longer reads that score at all — enforced by a dedicated test, not just a comment.

Of the baseline's **8 MISSING** requirements, **all 8 are now MET**. Of its **6 PARTIAL** requirements, **4 are now MET**; the remaining 2 (freshness display, value-driver tagging) were not touched today and are exactly as they were this morning — not attempted, not regressed. The three-way open question the baseline flagged — persona, geography and business-domain filtering, *"your call, not a coding decision"* — was answered by building all three.



None of this was free, and it should not be presented as finished — three things are worth saying up front, not discovered later:

- **Right-to-win / CRM data still does not exist anywhere in the app.** The new "why this matters" field and the Presales sort both have real, tested code paths waiting for it — today they honestly say so instead of inventing a number.
- **The new Next/Later independence is real but the corpus is thin.** 2 of 21 spaces land in "Next" today (9.5%) — the rule works, the live evidence base just isn't deep enough yet for more of it to show. Section 3 shows why with a real example.
- **A handful of countries — Greece foremost — don't yet resolve to a geography region.** They are reported as gaps, not silently misfiled into the wrong bucket. See Section 4.

Section 2 walks the baseline's own scorecard, row for row, with what changed. Sections 3–5 go deeper on the score, the new dimensions, and the new workflow layer. Section 6 is what's left for tomorrow.

Updated scorecard

Same rows as the baseline document. "Status" is where it stands now.

REQUIREMENT	STATUS THIS MORNING	WHAT CHANGED TODAY	STATUS
MVP technical scope — slide 19 <i>The briefing's own must-have bar.</i>			
An attractiveness scoring system	MISSING — the ranking number was a raw volume ratio	Five weighted components (market signal strength 30%, source credibility 20%, evidence quality 20%, novelty & momentum 15%, strategic relevance 15%), each independently computed and shown. A component with insufficient data is excluded and the rest rescaled — never silently zeroed. <code>services/attractiveness.py — combine()</code>	MET
Scoring must explain itself — slide 16 <i>Stated as a hard requirement, independent of any formula.</i>			
The model must explain the number, not only produce it	MISSING — one generated sentence, no breakdown	The detail page now carries a full "Why this score" breakdown (all 5 components, their weight, their value), a "Why this timing" breakdown (the exact counts and rule that produced the Now/Next/Later badge), and a "Signal types behind it" panel naming the distinguishing question each signal answered. <code>pages/opportunity_detail.py · components/ui.py — attractiveness_row, horizon_check_row, signal_type_row</code>	MET
Attractiveness score present	PARTIAL — present, but measured volume	Present and measures what it claims to.	MET
Urgency / time horizon	PARTIAL — a threshold cut on the volume proxy; "Next" was structurally unreachable (0 of 21 spaces)	Fully independent rule: signal type + signal date + source count decide the badge. "Now" requires ≥2 concrete signals from ≥2 distinct sources, one recent; a single tender lands in "Next", not "Now". A dedicated test asserts the calculation never reads the attractiveness score. Live result: 0 Now / 2 Next / 19 Later — see Section 3 for why "Next" is still thin. <code>services/horizon.py · tests/test_signal_horizon.py</code>	MET

Updated scorecard, continued

Requirement	Status this morning	What changed today	Status
Signal type	MISSING — no field existed	Six mutually-exclusive types (buying signal, regulation, proof signal, competitor move, market trend, tech maturity), each defined by one distinguishing question embedded in the classification prompt, with a fixed tie-break order. Deterministic for TED/OCDS/CORDIS, LLM-classified for RSS. 34 of 34 live classifications now carry a type. Pipelineteamfile/common/signal_types.py · collector/signal_route.py	MET
"Why hot now", "why this matters", recommended next action as distinct fields	MISSING — one hardcoded constant, identical on every topic	Three independently composed, fully deterministic fields — no LLM call generates them. "Why hot now" is one clause per recent signal, strongest type first. "Why this matters" is a domain/vertical clause plus a right-to-win clause (honestly empty today — see Section 6). "Recommended move" is a 9-cell role-mode × horizon matrix, so a Strategist and a Salesperson get different text for the same topic. services/explanations.py	MET
User acceptance criteria — slide 13			
Trust	PARTIAL — sources yes, explanation no	Both present.	MET
Freshness	PARTIAL — computed, never rendered	Not touched today. Still computed in team_repository, still not read by any page.	PARTIAL
Filtering by vertical	MET	Unchanged.	MET

Updated scorecard, continued

REQUIREMENT	STATUS THIS MORNING	WHAT CHANGED TODAY	STATUS
Filtering by geography	MISSING — geography was a free-text company-profile attribute only	Deterministic from structured sources (TED buyer country, CORDIS consortium participant countries, OCDS buyer country), LLM-inferred from text for RSS. 14-region taxonomy matching how the business actually segments Europe. Multi-select chip filter on the Opportunities page, OR within, AND across dimensions. Pipelineteamfile/common/geography.py · services/geography.py	MET
Filtering by business domain	MISSING — no field anywhere	The 6 Orange domains (OX: Smart Industries, Connectivity, Cybersecurity, Cloud, CX, EX), derived deterministically from each space's technology and use case — no LLM call. Multi-select chip filter. Pipelineteamfile/common/business_domains.py	MET
Filtering by persona	MISSING — zero matches for "persona" in the codebase	8 buyer personas, weighted (not boolean) via a use-case × persona table unioned with a domain × persona overlay, plus a persona × vertical suppression list. Feeds both filtering and ranking (a dampened score multiplier). Selected via a single-select picker in Sales mode rather than the domain/geography chip pattern — a UI inconsistency worth smoothing over, not a missing capability. Pipelineteamfile/common/personas.py · services/personas.py	MET
Topic detail depth — slides 5 and 14			
Dated, typed signals mixing qualitative and quantitative evidence	PARTIAL — dated yes, typed no	Now typed too. See Signal type row above.	MET
Use cases and value drivers with value tags	PARTIAL — use case named, no value drivers or tags	Not touched today.	PARTIAL

Updated scorecard, continued

REQUIREMENT	STATUS THIS MORNING	WHAT CHANGED TODAY	STATUS
Numeric score breakdown from sub-criteria bars	MISSING — two numbers, no sub-criteria	Two real breakdown panels with progress bars: the 5-component Attractiveness breakdown, and the horizon rule's own counts (signals pointing at Now / Next / Later, distinct sources, which rule fired). components/ui.py — attractiveness_row, horizon_check_row	MET
Right-to-win: capabilities, proof points, partner ecosystem	NOT REQUIRED — slide 18, explicitly next-phase	Still not built, correctly. The new "why this matters" field and the detail page's "Right to win & proof points" region are both wired and tested against this data the moment it exists — today they render an honest absence, not a placeholder number.	NOT REQUIRED
External signal taxonomy (slide 19, extra content)	NOT REQUIRED, but flagged as <i>"the enabler for the slide-16 signal-typing gap — worth deciding on together"</i>	Built. This is exactly the Signal type work above — the item the baseline report suggested as worth discussing turned out to be the same enabling work the slide-16 fix needed anyway.	MET

The baseline's open question, answered. "The briefing contradicts itself on persona and geography filtering... it is your call, not a coding decision, and it should be made before anyone starts." The call made today was to build all three (domain, persona, geography) rather than wait — they are now real, filterable, tested dimensions. If the MVP slide's narrower scope was actually intended, that's worth confirming, but nothing here needs to be undone either way; all three degrade honestly when their inputs are thin (see Section 4).

Deep dive: the Attractiveness score, now real

What the baseline report's Section 3 showed as broken, and what replaced it.

SPECIFIED (slide 17, illustrative weights)

- **30% Market signal strength** — how visible across external sources
- **20% Source diversity** — independent-source coverage
- **20% Evidence quality** — is the signal relevant and credible
- **15% Novelty and momentum** — is the topic emerging or rising
- **15% Strategic relevance** — does it fit Orange Business

IMPLEMENTED TODAY (services/attractiveness.py)

- **30% Market signal strength** — recency-decayed evidence density, ~9-month half-life, normalised against the strongest space this run
- **20% Source credibility** — mean trust score from a 9-tier publisher-category anchor table (renamed from "diversity"; rejected as the wrong signal given how sources are found)
- **20% Evidence quality** — mean classifier confidence, now under an anchored rubric instead of an unanchored self-report
- **15% Novelty & momentum** — percentile rank of recent-vs-prior growth against every other space this run, not a fixed curve (most spaces have 1–2 articles; a fixed curve was pure noise)
- **15% Strategic relevance** — match against priorities set on the Company page

```
combined_score =  $\Sigma(\text{weight} \times \text{value}) / \Sigma(\text{weight of available components})$ 
  - components with no data are excluded, remaining weights rescale, never counted as zero
```

One bug found and fixed along the way

`seed_hardcoded_sources()`'s upsert set the trust category but not the audited-at timestamp, so TED and CORDIS — the two highest-trust feeds, backing 18 of 34 classifications — stayed flagged "unaudited" and silently zeroed the Source credibility component on almost every space. Found and fixed before this report was written; verified against live data afterward.

One real topic, in the app right now

Opportunity space #28 as the interface presents it today.

Aerospace · Environmental/emissions monitoring × Machine Learning

opportunity_spaces id = 28 · 1 linked article · domain: Cloud · region: Benelux

ATTRACTIVENESS

75 / 100

3 of 5 components
active, 1 excluded

EVIDENCE CONFIDENCE

86%

mean classifier
confidence

HORIZON

Later

evidenced only by
tech_maturity

MOMENTUM

New

no earlier evidence to
compare

The breakdown the interface shows for the score above:

COMPONENT	WEIGHT	VALUE	STATUS
Market signal strength	30%	62 / 100	AVAILABLE
Source credibility	20%	100 / 100	AVAILABLE
Evidence quality	20%	86 / 100	AVAILABLE
Novelty & momentum	15%	50 / 100	AVAILABLE
Strategic relevance	15%	—	NO DATA YET

Strategic relevance is genuinely unavailable here — no priorities are currently set on the Company page (they were cleared during this session's own testing). This is the rescaling rule in action: the 75/100 score above is computed from the other four components only, not diluted by a fabricated zero.

What the interface says about why:

Why hot now: "Tech maturity: CORDIS project status SIGNED – funded research still..."

Why this matters: "A Cloud opportunity for Aerospace — no direct right-to-win evidence yet — early-stage watch."

Recommended move (Strategist mode): "Add to the watchlist for the safety manager — keep the pitch exploratory, evidence is still building."

Three sentences, three different composition rules, zero LLM calls to produce them — see Section 5.

The horizon fix

The baseline report's chart showed "Next" was structurally unreachable. Here is the same chart, and why it's still thin rather than declared solved.

This morning: horizon = "Now" if relevance >= 82 else "Next" if relevance >= 68 else "Later"

– a threshold cut on the same volume-proxy score. Largest space held 2 articles, so relevance could only be 50 or 100. Nothing could ever land between.

Now: horizon = f(signal_type, signal_date, event_date, source_count) – never reads the score.

– asserted by a test: a high-volume space evidenced only by tech_maturity signals must return Later; a low-volume space with one recent tender must return Next.

THIS MORNING

0 / 21

spaces could ever reach "Next"

NOW

2 / 21

spaces genuinely in "Next" (9.5%)

Why still only 2, honestly

25 of the 34 live classified signals are CORDIS tech_maturity records — funded research that hasn't reached a buyer yet, which correctly lands in "Later". The rule is not the bottleneck; the live corpus' composition is. One of the two "Next" spaces:

Healthcare · Diagnostic support × Machine Learning — 1 proof signal + 1 tech-maturity signal, but only 1 of 2 sources and 1 of 2 concrete signals needed for "Now" — the interface literally states: *"concrete evidence exists but does not converge: only 1 of 2 concrete signals, only 1 of 2 distinct sources."* That sentence is the rule explaining itself on a real space, not a summary written after the fact.

Recommendation: re-measure this distribution after the next real collection run rather than loosening the convergence thresholds now — narrowing them to force more "Next" results would defeat the reason this was rebuilt.

Three new dimensions, and their honest gaps

All three were MISSING in the baseline report. Each is real, tested, and each has a reported — not hidden — limitation.

Business domain

6 domains, derived deterministically as the union of a space's technology-domain mapping and its use-case-domain mapping (technology always wins for the single-value "primary" badge). Live distribution: OX: Smart Industries 11, Cloud 6, CX 2, Cybersecurity 2. No domain fell under the 5% coverage floor.

`Pipelineteamfile/common/business_domains.py`

Target persona

8 buyer personas (CIO, CISO, CDO, COO, CMO/CX, IT & Network, Quality Manager, Industrial Safety Manager), weighted Primary/Secondary/Peripheral rather than boolean — the weight is what makes ranking possible, not just filtering. All 21 live spaces resolve at least one persona. A persona × vertical suppression list zeroes out implausible pairs (e.g. Industrial Safety Manager × Retail).

`Pipelineteamfile/common/personas.py`

Gap: persona selection is currently a single-select picker inside Sales mode, not a standalone multi-select chip filter like domain and geography — same underlying capability, different UI treatment. Worth unifying.

Geography

14 regions matching how the business actually segments Europe (France and Germany each standalone, DACH meaning Switzerland+Austria only — deliberate, not a mistake to "fix"). Deterministic for TED/CORDIS/OCDS, LLM-inferred for RSS. 16 of 21 spaces resolve a region; 5 render honestly as "Global / unspecified" rather than being excluded.

`Pipelineteamfile/common/geography.py`

Reported gap, not silently misfiled: Greece appears in 12 of the 21 live spaces and has no home in the region table — it doesn't fall under Southern Europe (Italy/Spain/Portugal/Israel only), and the code deliberately refuses to push it into "Asia" via a naive continent fallback or bury it in "Global". It returns unresolved and is flagged. Cyprus, Malta, Turkey, Russia, the Western Balkans, Moldova and Belarus sit in the same hole. Fix is a one-line taxonomy edit plus a backfill rerun, once someone decides where Greece belongs.

A new workflow layer: Role Mode

Not in the baseline scorecard at all — a new capability, not a gap fix.

Three view modes — **Strategist**, **Sales**, **Presales** — implemented as a pure view configuration, never written to the opportunity record. Switching mode changes three things at once: the default filters (Sales asks for a persona before showing results), the sort key (Sales sorts by a persona-weighted score — base score × a dampened multiplier that reorders without letting a weak-fit topic outrank a strong one), and which parts of the detail page lead, sit standard, or start collapsed. No region is ever fully hidden in any mode — asserted by a test rendering all three modes and checking every region is still in the page.

Mode only re-seeds filters if the user hasn't touched them — changing a filter is a normal action, never silently discarded by switching mode. Sales mode also switches the list view from a card grid to one wide card per topic, since the brief describes it being used live in front of a customer.

`services/role_modes.py` · `config/role_modes.json`

Deterministic explanation fields

The baseline report's "Recommended move" — a single hardcoded constant, identical on every topic — is replaced by three composed fields built from typed clauses, not sentence templates with slots:

- **Why hot now** — one clause per recent qualifying signal, capped at 3, strongest signal type first.
- **Why this matters** — always exactly 2 clauses: a domain/vertical framing sentence, then a right-to-win sentence (today, always the honest fallback — see Section 6).
- **Recommended move** — a 9-cell role-mode × horizon matrix plus an action clause keyed to the dominant signal type; genuinely different text per mode for the same topic, closing the exact bug the baseline report called out by file and line number.

No structured "entity" field exists on a signal record (only free-text rationale) — clauses lean on a type lead-in plus a truncated rationale rather than guessing at extraction. Deliberately rejected using the article's source name as a stand-in actor (e.g. "Defense Daily committed...") since that names the publisher, not the actor, which would read as confidently false.

What's left for tomorrow

Ordered by value returned per unit of work, same convention as the baseline report.

1 **Decide where Greece (and its neighbours) belong in the region table**

SMALL · A DECISION, THEN A ONE-LINE CONFIG EDIT + BACKFILL

12 of 21 live spaces touch Greece and it currently resolves to nothing. Also affects Cyprus, Malta, Turkey, Russia, the Western Balkans, Moldova, Belarus.

2 **Unify the persona filter into the domain/geography chip pattern**

SMALL · UI CONSISTENCY, NO NEW LOGIC

The weighting and threshold logic already exists and is tested — this is purely presentational.

3 **Close the client_relevance dead path from the baseline report**

SMALL · ONE LINE OF PROMPT, OR REMOVE THE COLUMNS

Still unaddressed — the baseline report's recommendation 4. Strategic relevance now works through a different, real mechanism (Company-page priorities), which makes the old dead columns even more worth removing rather than leaving as a trap for a future reader.

4 **Right-to-win / CRM data — the one remaining structural gap**

MEDIUM/LARGE · NEEDS A DATA SOURCE DECISION FIRST

Affects "why this matters", the Presales sort (still falls back to Attractiveness with a visible note), and the "Right to win & proof points" region. Every consuming code path is already built and tested against synthetic data — this is a data-sourcing decision, not further engineering design.

5 **Render "last updated" and decide on value-driver tagging**

SMALL EACH · UNTOUCHED SINCE THE BASELINE REPORT

Freshness data is computed but not shown anywhere; value drivers/tags were never part of today's scope.

Nothing described in this report has been committed. 54 changed or new files sit in the working tree, all verified against 207 passing automated tests and the running local app before this document was written.